

# HPC ClusterLevel Validation & Performance Benchmarking Superbench, NCCL, and Linpack Methodology

By: M. EL AZZOUZI

**Scope:** This section defines the **clusterwide validation and benchmarking phase** after individual node acceptance. It ensures the cluster as a whole meets or exceeds design targets in GPU, network, and systemlevel response performance.

## 1) Objectives

- Verify **GPU fabric coherence** across the cluster (intranode and internode).
- Validate **NCCL, RDMA, and IB fabric health** at scale.
- Benchmark **endtoend performance** using **SuperBench** (synthetic + deep learning workloads).
- Detect topology or imbalance issues (e.g., link skew, GPUtoGPU latency asymmetry, NUMA misbinding).
- Generate an **automated health & performance heatmap** for management and future regression comparison.
- (Optional) Conduct **Top500/HPL benchmark** to rank the cluster performance against public metrics.


## 2) Tools and Frameworks

Component	Purpose	Notes
<b>NCCL Tests</b>	Validate allreduce, alltoall, broadcast, and reduce operations over IB/NVLink	Included in nccl-tests container or SuperBench suite
<b>SuperBench</b>	Endtoend benchmark suite for AI/HPC systems	Supports GPU, CPU, and network tests; JSON output with visualization; official NVIDIA + Microsoft collaboration
<b>DCGM Exporter</b>	Health metrics (temps, clocks, power, ECC, XIDs)	Should be enabled clusterwide during tests
<b>ibdiagnet / ibqueryerrors</b>	IB topology & counter verification	Part of MOFED; run before and after performance tests
<b>Grafana / Prometheus dashboards</b>	Visualization and anomaly detection	Used to produce performance heatmaps and live monitoring

**HPL (Top500)**

Optional largescale Linpack  
benchmark

To measure theoretical peak TFLOPS and  
compare against rankings



## 3) PreValidation Steps

1. **Ensure all nodes are accepted** from the nodelevel playbook.
2. Confirm:
  - o Fabric Manager, DCGM, and IB interfaces are healthy on all GPU nodes.
  - o Kubernetes node labels reflect correct `gpu.count` and `ib.ports`.
  - o RDMA is enabled: `/dev/infiniband/uverbs*` visible inside GPU containers.
  - o NCCL P2P and socket networks are configured:
    - o `export NCCL_DEBUG=INFO`
    - o `export NCCL_IB_HCA=mlx5_0,mlx5_1`
    - o `export NCCL_SOCKET_IFNAME=eth0`
    - o `export NCCL_NET_GDR_LEVEL=2`
3. Sync all system clocks and NTP across nodes.
4. Record cluster topology:
  5. `ibnetdiscover > ib_topo.txt`
  6. `ibdiagnet -r > ib_route_health.txt`

## 4) Quick NCCL Fabric Sanity Test

Run a short distributed NCCL test job across representative nodes.

```
# 4node NCCL test example
kubectl create -f nccltest.yaml
```

### Sample `nccl-test.yaml`:

```
apiVersion: batch.volcano.sh/v1alpha1
kind: Job
metadata:
  name: ncclsanity
  namespace: gcrdiag
spec:
  queue: gcrdiag
  minAvailable: 4
  plugins:
    ssh: []
  tasks:
    - replicas: 4
      name: ncclsanity
      template:
        spec:
          containers:
            - name: nccl
              image: nvcr.io/nvidia/nccl-tests:2.23.4-cuda12.4
              command: ["/bin/bash", "-c"]
```

```
args:  
- >  
    mpirun -np 32 --allow-run-as-root -bind-to none \  
        ./build/all_reduce_perf -b 8 -e 1G -f 2 -g 8  
restartPolicy: Never
```

### Expectations:

- All ranks initialize within seconds.
- Throughput within  $\pm 10\%$  of lab baseline (e.g.,  $\geq 360$  GB/s nodetonode on HDR400).
- No NCCL warnings about topology mismatch or IB channel fallback.

**Outcome:** A quick confidence check that GPUs, IB fabric, and NCCL runtime are coherent.

## 5) SuperBench Cluster Validation

SuperBench provides **modular validation and benchmarking** for both singlenode and fullcluster runs.

### 5.1 Deployment

Install via Helm:

```
helm repo add superbench https://superbench.github.io/charts  
helm upgrade --install superbench superbench/superbench \  
    --namespace superbench --create-namespace \  
    --set cluster.mode=volcano --set benchmark.mode=distributed
```

### 5.2 Configuration

Edit `values.yaml` to specify tests and node pools:

```
superbench:  
  enable: true  
  mode: distributed  
  node_selector:  
    k8s.lambdai.ai/node-type: gpu  
  benchmarks:  
    - name: gpu-burn  
    - name: nccl-allreduce  
    - name: memory-bandwidth  
    - name: ib-loopback  
    - name: kernel-launch  
    - name: disk-bw  
  duration: 1800  
  output:  
    format: json
```

```
save_path: /results/superbench-$ (date +%F).json
```

### 5.3 Execution Modes


Mode	Description	Expected Runtime
<b>Halfcluster test</b>	Run on 50% of GPU nodes (balanced across racks)	~1 hour
<b>Fullcluster validation</b>	Run on all GPU nodes simultaneously	2–4 hours depending on scale

### 5.4 PostProcessing

```
superbench analyze result --input /results/*.json --output /results/summary.xlsx
```

- Generates **performance summaries per node** (latency, bandwidth, FLOPS, memory BW, etc.).
- Integrate results into **Grafana heatmap dashboard**:
  - Xaxis: node index or hostname
  - Yaxis: test type (NCCL, IB, GPU, CPU)
  - Color: relative performance (green  $\geq$  baseline, yellow = warning, red = regression)

**Deliverable:** A visual **cluster health/performance heatmap** + Excel summary stored under `/validation_reports/<date>/.`




## 6) Additional Network & System Tests

Test	Tool	Command / Example	Pass Criteria
<b>IB port error check</b>	ibqueryerrors	ibqueryerrors -v	Zero critical counters before/after test
<b>Fabric routing check</b>	ibdiagnet	ibdiagnet -r	No invalid or missing routes
<b>Latency per GPU pair</b>	nvsysdiag or nvlink_bandwidth	nvidia-smi nvlink -s	No disabled NVLinks
<b>Collective sync test</b>	alltoall_perf	mpirun -np <n>	Stable mean latency
<b>CPU/RAM imbalance</b>	numactl --hardware + stream	stream --benchmark	Balanced throughput per socket
<b>Filesystem IO</b>	fio	See nodelevel tests	IO bandwidth $\geq$ baseline

**Thermal drift**      `dcmi dmon`

During SuperBench run      Temp deltas < 10°C per node



## 7) Reporting Template

Section	Description	Data Source
Hardware Summary	GPU/CPU/RAM/IB per node	Inventory database
NCCL Quick Test	Sanity throughput (GB/s)	nccl-tests logs
SuperBench Summary	Pertest and pernode metrics	superbench analyze output
Heatmap	Visualized node health	Grafana or Excel heatmap
IB Topology	Connectivity map	ibdiagnet results
Regression Detection	$\Delta$ vs previous validation	SuperBench JSON diff
Anomalies	Xid, ECC, link errors	DCGM + ibqueryerrors logs
Recommendation	Accept / Rerun / Quarantine nodes	Engineer decision

## 8) Optional: Top500 / HPL Benchmark (If Time Allows)

Running a **High Performance Linpack (HPL)** test provides an approximate **Top500 ranking**.

## 8.1 Toolchain

- Container: `nvcr.io/nvidia/hpc-benchmarks:latest`
- Scheduler: run under Volcano with full GPU access.
- Config file: define problem size `N`, process grid `PxQ`, and block size `NB`.

## 8.2 Example Job

```
mpirun -np 256 --map-by ppr:8:node \
./xhpl
```

## 8.3 Output

- $R_{max}$  = measured TFLOPS
- $R_{peak}$  = theoretical TFLOPS
- Compute **Efficiency =  $R_{max}/R_{peak}$**  (target  $\geq 85\%$ )
- Compare against **Top500** list of similar hardware generations.

## 8.4 Usefulness

- Validates CPU–GPU–network orchestration at extreme scale.
- Serves as a singlenumber metric for management and marketing.
- Archive the run under `/benchmarks/top500/<date>/`.

# 9) Acceptance Criteria (Cluster Level)

Category	Pass Condition	Action on Fail
GPU/NCCL	$\geq 90\%$ of baseline throughput; all ranks initialize	Investigate failing nodes; rerun after cordon
IB Fabric	0 critical errors; all ports at full speed	Replace cable/transceiver; revalidate
SuperBench summary	$\geq 95\%$ green cells in heatmap	Analyze yellow/red nodes
System stability	No node crash, Xid, or fabric reset during test	Quarantine failing nodes
Optional HPL	$\geq 80\%$ efficiency vs Rpeak	Log and compare vs historical

# 10) Deliverables

1. **SuperBench Report Bundle:** JSON + XLSX summaries.
2. **Cluster Heatmap:** PNG or Grafana dashboard snapshot.
3. **NCCL & IB logs:** archived under `/validation_reports/<date>/logs/`.

4. Linpack Top500 report (if executed).
5. Executive Summary: key scores, bottlenecks, actions.

## 11) Future Enhancements

- Integrate automated NCCL & SuperBench validation into CI/CD for new nodes.
- Baseline heatmap autopublish in Grafana after each intake.
- Autodetect regressions (e.g., via zscore anomaly detection on throughput).
- Use collected data for **capacity planning** and **vendor performance SLAs**.

**End - Part 2**